

# Anomaly Detection with Robust Deep Autoencoders

original article by C. Zhou and R. C. Peffenroth

*Alessandro Trenta*

Scuola Normale Superiore

- 1 Background
- 2 Robust Deep Autoencoders
- 3 RDAE training

# Deep Autoencoders

- A Deep Autoencoder (DAE) is constituted by two main components: an encoder  $E$  and a Decoder  $D$ .
- The main objective of a DAE is to learn the identity map so that the reconstruction  $\bar{X} = D(E(X))$  is as close as possible to the original input  $X$ .
- The encoder and decoder functions  $E, D$  can be any kind of mapping between the data space and the coding space. Usually they are Deep Neural Networks e.g. a feed forward network or even more complex models such as Long Short Ter Memory (LSTM).
- The objective is usually to find the minimum reconstruction error w.r.t. some parametrized encoding and decoding functions and a distance (in this case the  $L_2$  norm)

$$\min_{\theta, \phi} \|X - D_{\theta}(E_{\phi}(X))\|_2 \quad (1)$$

# Principal Component Analysis

- Assume to have a set of  $N$  samples of  $n$  dimensional data, so that  $X \in \mathbb{R}^{N \times n}$  s.t. each column has 0 mean (we can just shift the data to fulfill this request).
- Principal Component Analysis (PCA) is defined as an orthogonal linear transformation such that the new coordinate system of  $\mathbb{R}^n$  satisfies: the  $i$ -th component of the coordinate system has the  $i$ -th greatest data variance if we project all samples on that component.
- Ideally we are trying to fit a  $n$ -ellipsoid into the data. The length of an axis of the ellipsoid represents the variance of data along that axis.
- PCA is often used for dimensionality reduction or encoding: we can project the data on the first  $k < n$  principal components.

# Principal Component Analysis

Mathematically we can define:

$$w_1 = \arg \max_{\|w\|_2=1} \|Xw\|_2^2 = \arg \max_w \frac{w^T X^T X w}{w^T w} \quad (2)$$

for the first component. Then for the  $k$ -th component we first subtract the first  $k-1$  principal component from  $X$

$$\hat{X}_k = X - \sum_{i=1}^{k-1} X w_i w_i^T \quad (3)$$

and finally solving again the similar problem:

$$w_k = \arg \max_{\|w\|_2=1} \|\hat{X}_k w\|_2^2 = \arg \max_w \frac{w^T \hat{X}_k^T \hat{X}_k w}{w^T w} \quad (4)$$

# Robust Principal Component Analysis

- Robust Principal Component Analysis (RPCA) is a generalization of PCA that aims to reduce the sensitivity of PCA to outliers.
- The idea is to find a low-dimensional representation of data cleaned from the sparse outliers that can disturb the PCA process.
- We therefore assume that data  $X$  can be represented as  $X = L + S$ :  $L$  has low rank and is the low-dimensional representation of  $X$  while  $S$  is a sparse matrix consisting of the outlier elements that cannot be captured by the representation.

# Robust Principal Component Analysis

- The problem can be addressed as:

$$\min_{L,S} \rho(L) + \lambda \|S\|_0 \quad (5)$$

$$\text{s. t. } \|X - L - S\|_F^2 = 0 \quad (6)$$

where  $\rho(\cdot)$  is the rank of a matrix and we used the zero norm.

- This optimization problem is NP-hard and tractable only for small matrices.
- Usually it is substituted by the following problem, which is convex and tractable also for large matrices:

$$\min_{L,S} \|L\|_* + \lambda \|S\|_1 \quad (7)$$

$$\text{s. t. } \|X - L - S\|_F^2 = 0 \quad (8)$$

where  $\|\cdot\|_*$  is the nuclear norm i. e. the sum of singular values of a matrix.

# Robust Deep Autoencoders

- The main idea behind Robust Deep Autoencoders (RDAE) is to combine the representation learning of DAEs and the anomaly detection capability of RPCA.
- Noise and outliers are incompressible in the lower dimensional space we want to represent our data in.
- The objective is to learn a good low dimensional representation except for few exceptions.
- We will see two RDAE typed, one for  $l_1$  regularization and one for  $l_{2,1}$ .



# RDAE with $l_1$ regularization

- The RDAE objective is to decompose data  $X = L_D + S$  just as in RPCA.
- By removing the noise  $S$  the autoencoder can better reconstruct  $L_D$ .
- As before, the best choice to obtain a sparse  $S$  would be to use a loss of the type  $\|S\|_0$  which counts the non-zero entries, solving the problem

$$\min_{\theta} \|L_D - D_{\theta}(E_{\theta}(L_D))\|_2 + \lambda \|S\|_0 \quad (9)$$

$$\text{s.t. } X - L_D - S = 0 \quad (10)$$

- The parameter  $\lambda$  controls the sparsity of  $S$  and plays an essential role.

# The role of $\lambda$

- As we said, the role of  $\lambda$  is very important.
- A smaller  $\lambda$  means that the norm of  $S$  plays a less important role and much of the loss will come from the DAE.
- The model will reconstruct better but recognize less outliers. This could be helpful if we want a more faithful representation e.g. for supervised tasks.
- A larger  $\lambda$ , instead, gives more importance to the norm of  $S$  as a loss.
- This means that the model will recognize more (or even too much) outliers, sacrificing some reconstruction performance.

# The true objective

- As for the RPCA the previous loss is non tractable. We then instead focus on the following problem:

$$\min_{\theta} \|L_D - D_{\theta}(E_{\theta}(L_D))\|_2 + \lambda \|S\|_1 \quad (11)$$


$$\text{s.t. } X - L_D - S = 0 \quad (12)$$

Notice two things:

- The autoencoder is trained with  $L_D$ , the part of its decomposition assumed to be outliers and noise free.
- There is no specific requirement about the DAE, in fact  $D_{\theta}$  and  $E_{\theta}$  are generic decoder, encoder functions.

## RDAE with $l_{2,1}$ regularization

- The RDAE with  $l_1$  penalization assumes that outliers and noise are not structured. The  $l_1$  penalty is indeed just a regularization to induce sparsity.
- In general we can have multiple types of errors: an instrument that corrupts an input feature or outliers where the input data is in some way structurally different from normal data.



# The $l_{2,1}$ norm

- The  $l_{2,1}$  norm is defined as ( $X \in \mathbb{R}^{N \times n}$ ):

$$\|X\|_{2,1} = \sum_{j=1}^n \|X_j\|_2 = \sum_{j=1}^n \left( \sum_{i=1}^N |X_{ij}|^2 \right)^{\frac{1}{2}} \quad (13)$$

- The  $l_{2,1}$  norm can be seen as introducing a  $l_2$  norm regularization over each feature and then adding a  $l_1$  regularization accross features.
- We can also do the other way around: to recognize data anomalies (by row) just apply the  $l_{2,1}$  norm to  $X^T$ .

- The final optimization problem for the RDAE with  $l_{2,1}$  regularization for data anomalies is then

$$\min_{\theta} \|L_D - D_{\theta}(E_{\theta}(L_D))\|_2 + \lambda \|S^T\|_{2,1} \quad (14)$$

$$\text{s.t. } X - L_D - S = 0 \quad (15)$$

- For detecting feature anomalies we just need to change the objective to

$$\min_{\theta} \|L_D - D_{\theta}(E_{\theta}(L_D))\|_2 + \lambda \|S\|_{2,1} \quad (16)$$

$$\text{s.t. } X - L_D - S = 0 \quad (17)$$

# The proximal operator

- To see in detail the training procedure for the RDAE we first need to consider the proximal operator.
- For general optimization problems of the form  $\min f(x) + \lambda g(x)$  where  $g$  is convex some of the most used methods require to find

$$\text{prox}_{\lambda, g}(x) = \arg \min_y g(y) + \frac{1}{2\lambda} \|x - y\|_2^2 \quad (18)$$

- In this case we then want to obtain a solution of the problems

$$\text{prox}_{\lambda, l_1}(x) = \arg \min_y l_1(y) + \frac{1}{2\lambda} \|x - y\|_2^2 \quad (19)$$

$$\text{prox}_{\lambda, l_{2,1}}(x) = \arg \min_y l_{2,1}(y) + \frac{1}{2\lambda} \|x - y\|_2^2 \quad (20)$$



- For the  $l_1$  norm, the solution to the proximal problem is

$$\text{prox}_{\lambda, l_1}(x) = \begin{cases} x_i - \lambda, & x_i > \lambda \\ x_i + \lambda, & x_i < -\lambda \\ 0, & x_i \in [-\lambda, \lambda] \end{cases} \quad (21)$$

which in the case of  $S \in \mathbb{R}^{N \times n}$  gets applied element by element.

- For the  $l_{2,1}$  norm, we obtain (let  $S_{\cdot j}$  be the column vector  $S_{ij}, j = 1, \dots, N$ )

$$(\text{prox}_{\lambda, l_{2,1}}(S))_{ij} = \begin{cases} S_{ij} - S_{ij} - \lambda \frac{S_{ij}}{\|S_{\cdot j}\|_2}, & \|S_{\cdot j}\|_2 > \lambda \\ 0, & \|S_{\cdot j}\|_2 \leq \lambda \end{cases} \quad (22)$$

if we are considering feature wise anomalies, substitute  $S$  with  $S^T$  for data anomalies.

# The main algorithm

- The method used to train the RDAE is the Alternating Direction Method of Multipliers (ADMM).
- The main idea is to optimize the problem

$$\min_{\theta} \|L_D - D_{\theta}(E_{\theta}(L_D))\|_2 + \lambda \|S^T\|_{2,1} \quad (23)$$

$$\text{s.t. } X - L_D - S = 0 \quad (24)$$

by doing it in two steps at each iteration.

- First, we fix  $S$  and optimize the DAE loss  $\|L_D - D_{\theta}(E_{\theta}(L_D))\|_2$  with backpropagation as usual.
- Then, we fix  $L_D$  and optimize the regularization term with the proximal method.





The full procedure is the following: given input  $X \in \mathbb{R}^{N \times n}$ , initialize  $L_D \in \mathbb{R}^{N \times n}$ ,  $S \in \mathbb{R}^{N \times n}$  as zero matrices,  $L_S = X$  and initialize the DAE randomly. For each iteration do:

- $L_D = X - S$
- Minimize  $\|L_D - D_\theta(E_\theta(L_D))\|_2$  with backpropagation.
- Set  $L_D = D(E(L_D))$  as the reconstruction.
- Set  $S = X - L_D$ .
- Optimize  $S$  using a  $\text{prox}_{\lambda, l}$  function of choice.
- If  $c_1 = \frac{\|X - L_D - S\|_2}{\|X\|_2} < \epsilon$  or  $sc_2 = \frac{\|LS - L_D - S\|_2}{\|X\|_2} < \epsilon$  we have early convergence.
- Set  $L_S = L_D + S$ .

Return  $L_D$  and  $S$ .

<https://github.com/AlexThirty/SaMLMfTSA>

Thank you!

-  Subutai Ahmad, Alexander Lavin, Scott Purdy, and Zuha Agha.  
Unsupervised real-time anomaly detection for streaming data.  
*Neurocomputing*, 262:134–147, 2017.  
Online Real-Time Learning Strategies for Data Streams.
-  Emmanuel J. Candes, Xiaodong Li, Yi Ma, and John Wright.  
Robust principal component analysis?, 2009.
-  Neal Parikh.  
Proximal algorithms.  
*Foundations and Trends in Optimization*, 1:127–239, 01 2014.
-  Chong Zhou and Randy C. Paffenroth.  
Anomaly detection with robust deep autoencoders.  
*Proceedings of the 23rd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, page 665–674, 2017.